

Time Series

Emerson Amirhosein Azarbakht

Time Series (TS)

Used in

- control of inventory, based on demand trends
- airline's decision to buy airplanes bc of passenger trends and decision to increase/maintain market share
- climate change decisions based on temperature change trends
- business/sales forecasting
- everyday operational decisions
- long-term effects of proposed water management policies by simulating daily rainfall and sea state time series
- understanding fluctuations in monthly sales
- basis for signal processing in telecommunications <?>
- disease incidence tracking, yearly rates
- census analysis
- tracking monthly unemployment rate; as an economic indicator used by decision makers

Used to

- to understand the past, and predict the future
- forecasting (predicting inference, a subset of statistical inference). assumes that present trends continue. This assumption cannot be checked empirically, but, when we identify the likely causes for a trend, we can justify the forecasting(extrapolating it) for a few time-steps at least
- anomaly detection
- clustering
- classification (assigning a time series pattern to a specific category: e.g. gesture recognition of hand movements in sign language videos)
- query by content ~ Content-based image retrieval

Data: a variable measured sequentially in time, or at a fixed [sampling] interval

serial dependence problem: observations close together in time tend to be correlated (serially dependent)

TS tries to explain this correlation (serial dependence) autocorrelation analysis examines this serial dependence <?>

conditions (assumptions of TS)

- Stationary process:
- if there's no systematic change in mean (no trend) AND
- if there's no systematic change in variance, AND
- if strictly periodic variations have been removed
 - i.e. the properties of one section of the data are much like those of any other section

- often we have a non-stationary TS => we need to remove trend and seasonality, to get a stationary residual, which then can be modeled using a stationary stochastic process
- Ergodic process: a stationary TS that we assume is sufficiently long time series that it characterises the hypothetical model. (~ independent of the starting point)

```
plot(AirPassengers)
start(AirPassengers)
```

```
## [1] 1949    1
```


```
end(AirPassengers)
```

```
## [1] 1960   12
```

```
frequency(AirPassengers)
```

```
## [1] 12
```

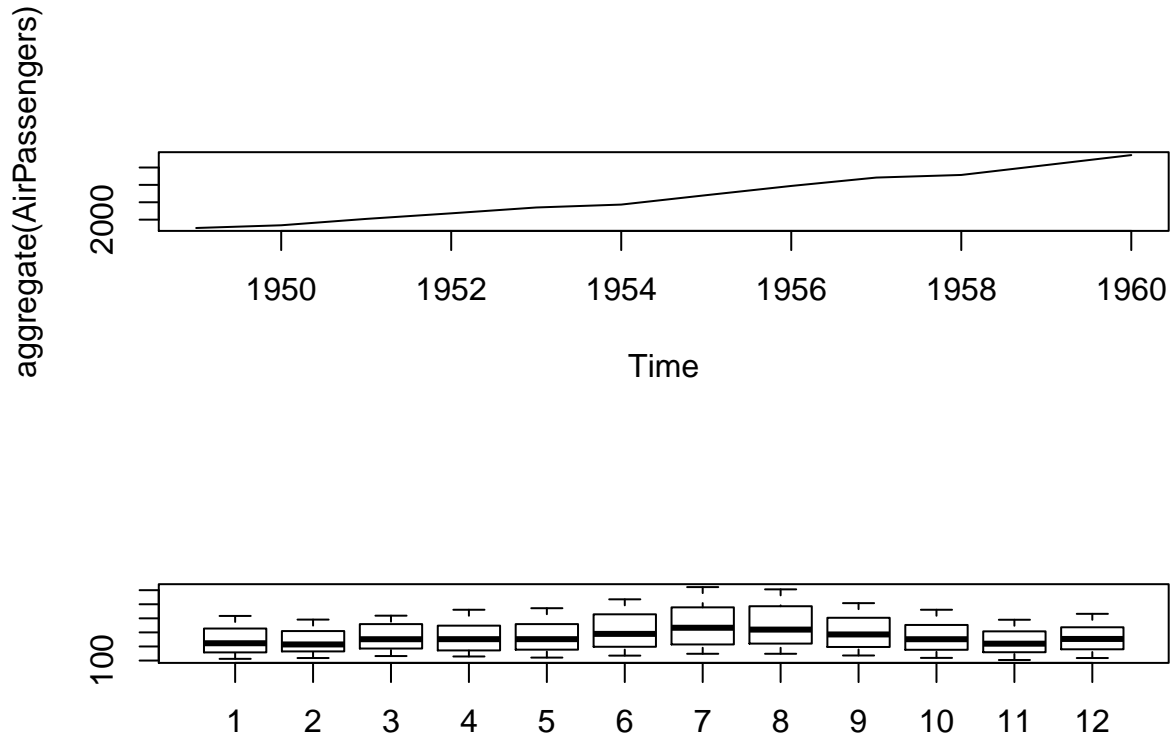
```
plot(AirPassengers)
```



```
summary(AirPassengers)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   104.0   180.0   265.5   280.3   360.5   622.0
```

```
layout(1:2)
# takes an input matrix for the location of each plot in the graphics window
plot(aggregate(AirPassengers))
boxplot(AirPassengers ~ cycle(AirPassengers))
```

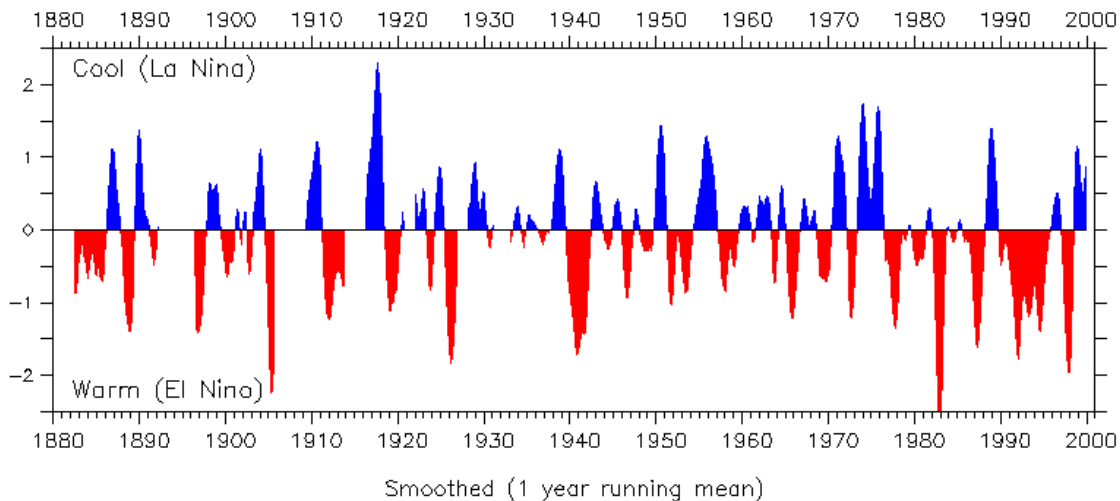


plotting shows *patterns*, and *features* of the data + *outliers* and *erroneous* values

patterns

1. trend = a non-periodic systematic change in a TS: a long-term change in the 'mean'
 - can be modeled simply by a linear increase or decrease. (only if it's deterministic ~ it's non-stochastic)
 - stochastic trend: seems to change direction at unpredictable times rather than displaying a consistent pattern (e.g. like the air passenger series); inexplicable changes in direction
2. seasonal variation = a repeating pattern within a fixed period (e.g. each year)
3. cycles = a non-fixed-period cycle (without a fixed period). example: El-Nino

Southern Oscillation Index



```
# monthly unemployment rate for the US state of Maine from January 1996 until August 2006
Maine.month <- read.table("http://staff.elena.aut.ac.nz/Paul-Cowpertwait/ts/Maine.dat", header = TRUE)
# header TRUE means treat first row as column names
attach(Maine.month)
str(Maine.month)
```

Monthly unemployment rate for a state

```
## 'data.frame': 128 obs. of 1 variable:
## $ unemploy: num 6.7 6.7 6.4 5.9 5.2 4.8 4.8 4 4.2 4.4 ...
```

```
head(Maine.month)
```

```
## unemploy
## 1 6.7
## 2 6.7
## 3 6.4
## 4 5.9
## 5 5.2
## 6 4.8
```

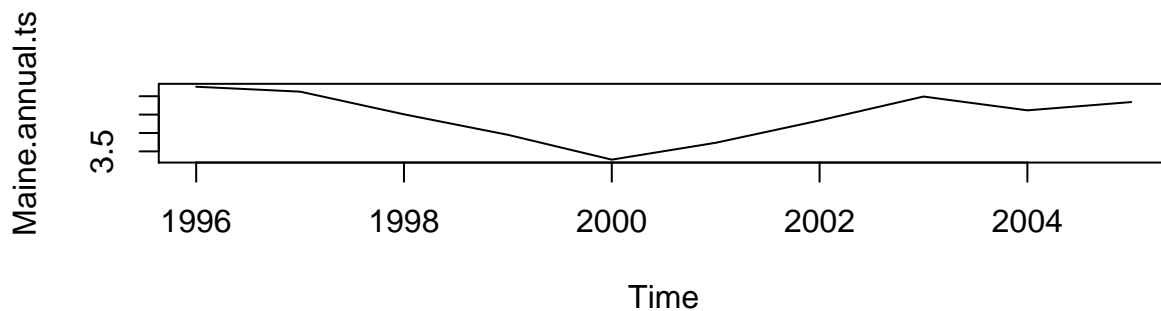

```
class(Maine.month)
```

```
## [1] "data.frame"
```

```
# it's a data.frame, not a ts object. So, we need to convert it to ts
Maine.month.ts <- ts(unemploy, start = c(1996, 1), freq = 12)
Maine.month.ts
```

```
##      Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
## 1996 6.7 6.7 6.4 5.9 5.2 4.8 4.8 4.0 4.2 4.4 5.0 5.0
## 1997 6.4 6.5 6.3 5.9 4.9 4.8 4.5 4.0 4.1 4.3 4.8 5.0
## 1998 6.2 5.7 5.6 4.6 4.0 4.2 4.1 3.6 3.7 4.1 4.3 4.0
## 1999 4.9 5.0 4.6 4.3 3.9 4.0 3.6 3.3 3.1 3.3 3.7 3.7
## 2000 4.4 4.4 4.1 3.5 3.1 3.0 2.8 2.5 2.6 2.8 3.1 3.0
## 2001 3.9 4.2 4.0 4.1 3.5 3.5 3.4 3.1 3.4 3.7 4.0 4.0
## 2002 5.0 4.9 5.0 4.7 4.0 4.2 4.0 3.6 3.7 3.9 4.5 4.6
## 2003 5.6 5.8 5.6 5.5 4.8 4.9 4.8 4.3 4.5 4.6 4.8 4.7
## 2004 5.6 5.6 5.5 4.8 4.2 4.3 4.2 3.8 4.0 4.2 4.6 4.6
## 2005 5.5 5.8 5.5 5.2 4.7 4.6 4.5 4.1 4.4 4.4 4.8 4.6
## 2006 5.3 5.6 4.9 4.6 4.2 4.4 4.4 3.9
```

```
Maine.annual.ts <- aggregate(Maine.month.ts)/12
layout(1:2)
plot(Maine.month.ts)
plot(Maine.annual.ts)
```



```
Maine.Feb <- window(Maine.month.ts, start = c(1996,2), freq = TRUE)
Maine.Aug <- window(Maine.month.ts, start = c(1996,8), freq = TRUE)
Feb.ratio <- mean(Maine.Feb) / mean(Maine.month.ts)
Aug.ratio <- mean(Maine.Aug) / mean(Maine.month.ts)
```

Multiple TS

```
ChocolateBeerElectricity <- read.table("http://staff.elena.aut.ac.nz/Paul-Cowpertwait/ts/cbe.dat", head=
class(ChocolateBeerElectricity)
```

```
## [1] "data.frame"
```

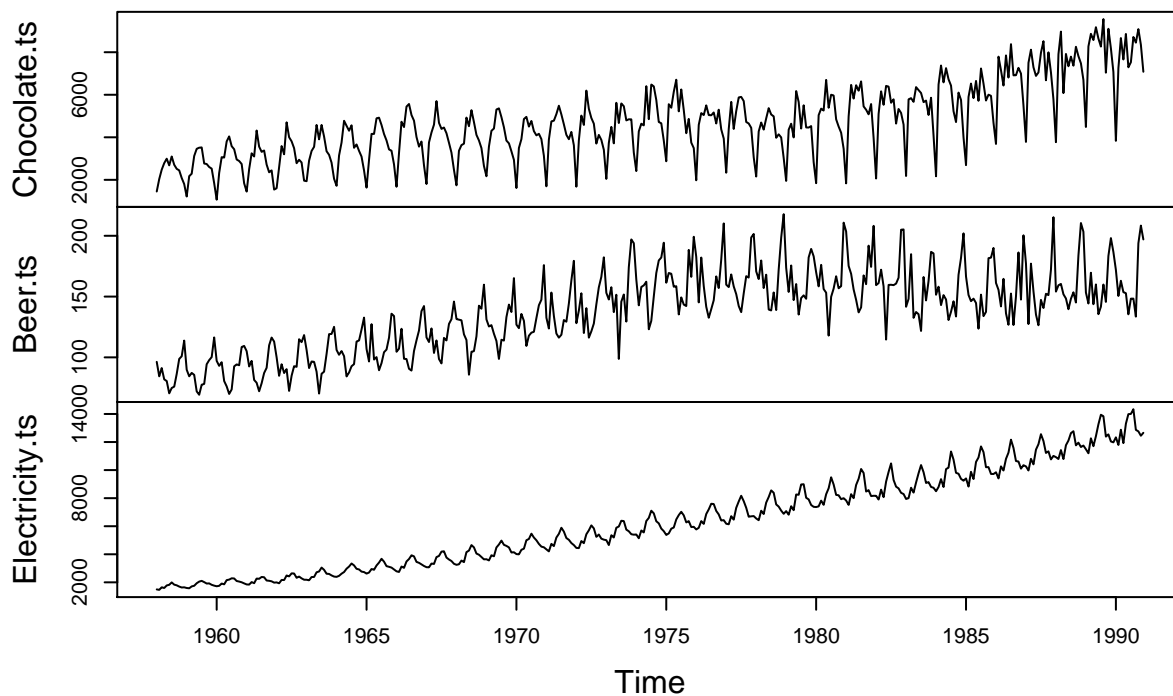
```
str(ChocolateBeerElectricity)
```

```
## 'data.frame': 396 obs. of 3 variables:
## $ choc: int 1451 2037 2477 2785 2994 2681 3098 2708 2517 2445 ...
## $ beer: num 96.3 84.4 91.2 81.9 80.5 70.4 74.8 75.9 86.3 98.7 ...
## $ elec: int 1497 1463 1648 1595 1777 1824 1994 1835 1787 1699 ...
```

```
Chocolate.ts <- ts(ChocolateBeerElectricity[,1], start = 1958, frequency = 12)
Beer.ts <- ts(ChocolateBeerElectricity[,2], start = 1958, frequency = 12)
Electricity.ts <- ts(ChocolateBeerElectricity[,3], start = 1958, frequency = 12)

plot(cbind(Chocolate.ts, Beer.ts, Electricity.ts))
```

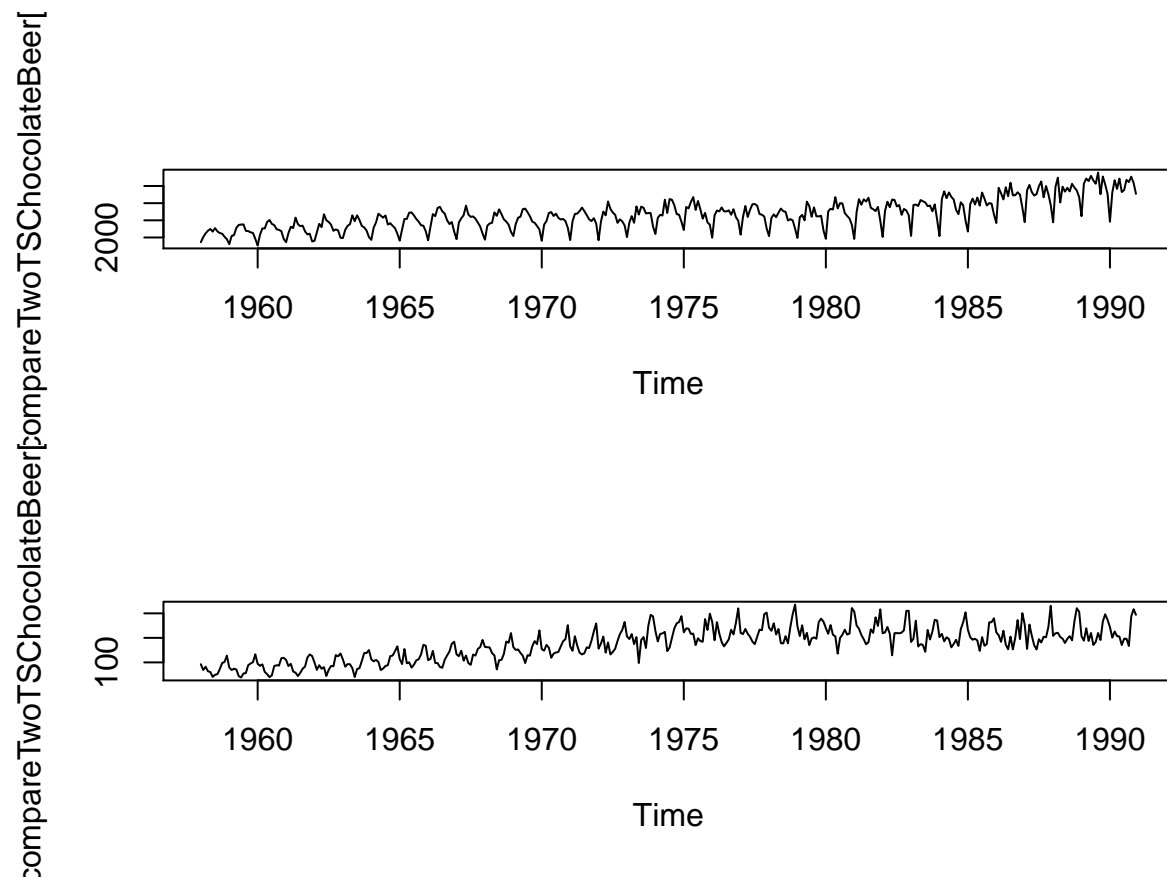
cbind(Chocolate.ts, Beer.ts, Electricity.ts)



```
# intersection of multiple TS
compareTwoTSChocolateBeer <- ts.intersect(Chocolate.ts, Beer.ts)
# bind time series which have a common frequency
?ts.intersect()
head(compareTwoTSChocolateBeer)
```

```
##      Chocolate.ts Beer.ts
## [1,]         1451    96.3
## [2,]         2037    84.4
## [3,]         2477    91.2
## [4,]         2785    81.9
## [5,]         2994    80.5
## [6,]         2681    70.4
```

```
layout(1:2)
plot(compareTwoTSSchocolateBeer[,1])
plot(compareTwoTSSchocolateBeer[,2])
```



```
cor(Chocolate.ts, Beer.ts)
```

```
## [1] 0.3774314
```

```
# correlation
```

```
# Global temperature anomalies from the monthly means over the period
GlobalTemperatures <- scan("http://staff.elena.aut.ac.nz/Paul-Cowpertwait/ts/global.dat")
```

```
# scan to read data into a vector or list from the console or file.
?scan
str(GlobalTemperatures)
```

Climate change

```
## num [1:1800] -0.384 -0.457 -0.673 -0.344 -0.311 -0.071 -0.246 -0.235 -0.38 -0.418 ...
```

```
GlobalTemperatures.ts <- ts(GlobalTemperatures, st = c(1856,1), end = c(2005,1), frequency = 12)
head(GlobalTemperatures.ts)
```

```
## [1] -0.384 -0.457 -0.673 -0.344 -0.311 -0.071
```

```
tail(GlobalTemperatures.ts)
```

```
## [1] 0.436 0.452 0.494 0.586 0.385 0.502
```

```
Global.annual <- aggregate(GlobalTemperatures.ts, FUN = mean)
head(Global.annual)
```

```
## [1] -0.3812500 -0.4611667 -0.4153333 -0.2252500 -0.3697500 -0.4003333
```

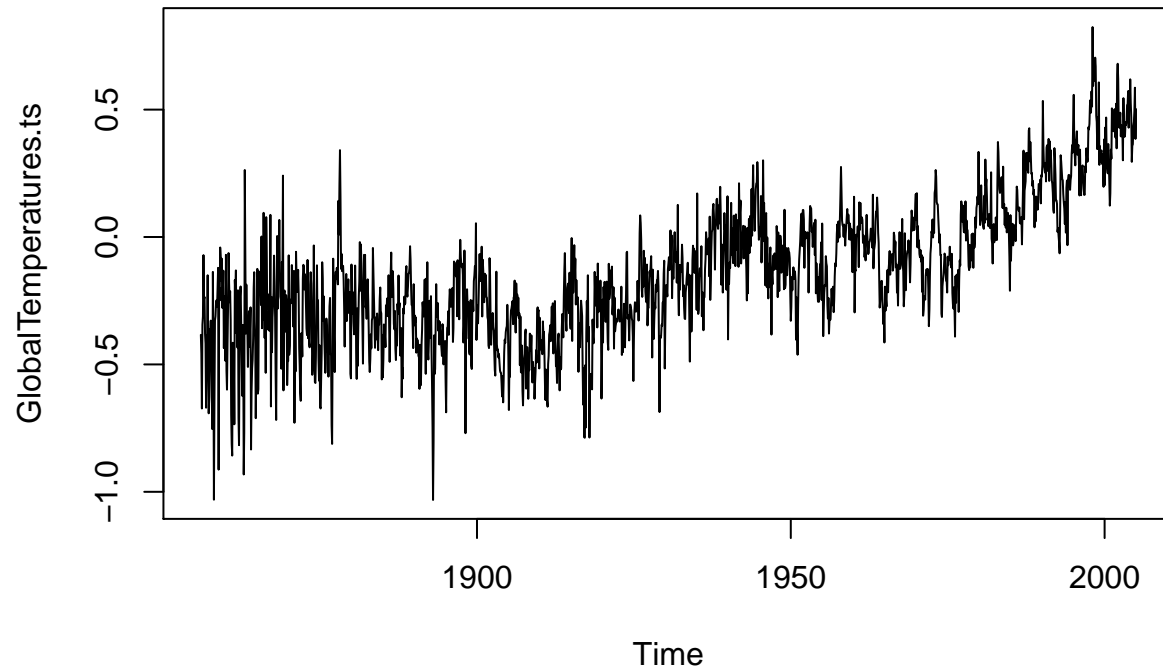
```
Global.annual
```

```
## Time Series:
## Start = 1856
## End = 2004
## Frequency = 1
## [1] -0.381250000 -0.461166667 -0.415333333 -0.225250000 -0.369750000
## [6] -0.400333333 -0.518916667 -0.273333333 -0.475333333 -0.262583333
## [11] -0.222416667 -0.289916667 -0.223583333 -0.305500000 -0.291333333
## [16] -0.339000000 -0.263083333 -0.329583333 -0.376500000 -0.421583333
## [21] -0.454750000 -0.212166667 -0.059500000 -0.288916667 -0.295166667
## [26] -0.245333333 -0.262416667 -0.317166667 -0.347583333 -0.349916667
## [31] -0.256083333 -0.345833333 -0.311916667 -0.202750000 -0.410416667
## [36] -0.351750000 -0.408583333 -0.447833333 -0.411083333 -0.362666667
## [41] -0.199083333 -0.184250000 -0.339416667 -0.252416667 -0.190916667
## [46] -0.257583333 -0.348583333 -0.446000000 -0.444166667 -0.370000000
## [51] -0.291916667 -0.505750000 -0.478750000 -0.448333333 -0.440500000
## [56] -0.461833333 -0.405916667 -0.393916667 -0.249333333 -0.161000000
## [61] -0.371416667 -0.498666667 -0.407333333 -0.289833333 -0.289750000
## [66] -0.215500000 -0.321333333 -0.295083333 -0.342250000 -0.244333333
## [71] -0.113416667 -0.217416667 -0.222916667 -0.354416667 -0.151833333
## [76] -0.096500000 -0.137500000 -0.238333333 -0.135666667 -0.170750000
## [81] -0.119416667 -0.023666667 0.073583333 -0.041916667 -0.081166667
## [86] 0.027250000 -0.017166667 -0.001833333 0.154083333 0.038583333
## [91] -0.114250000 -0.100666667 -0.092083333 -0.097666667 -0.207083333
## [96] -0.093000000 -0.020416667 0.043250000 -0.168416667 -0.189583333
## [101] -0.275166667 -0.004750000 0.064166667 0.013250000 -0.028416667
## [106] 0.014166667 0.002166667 0.034500000 -0.235583333 -0.167583333
```

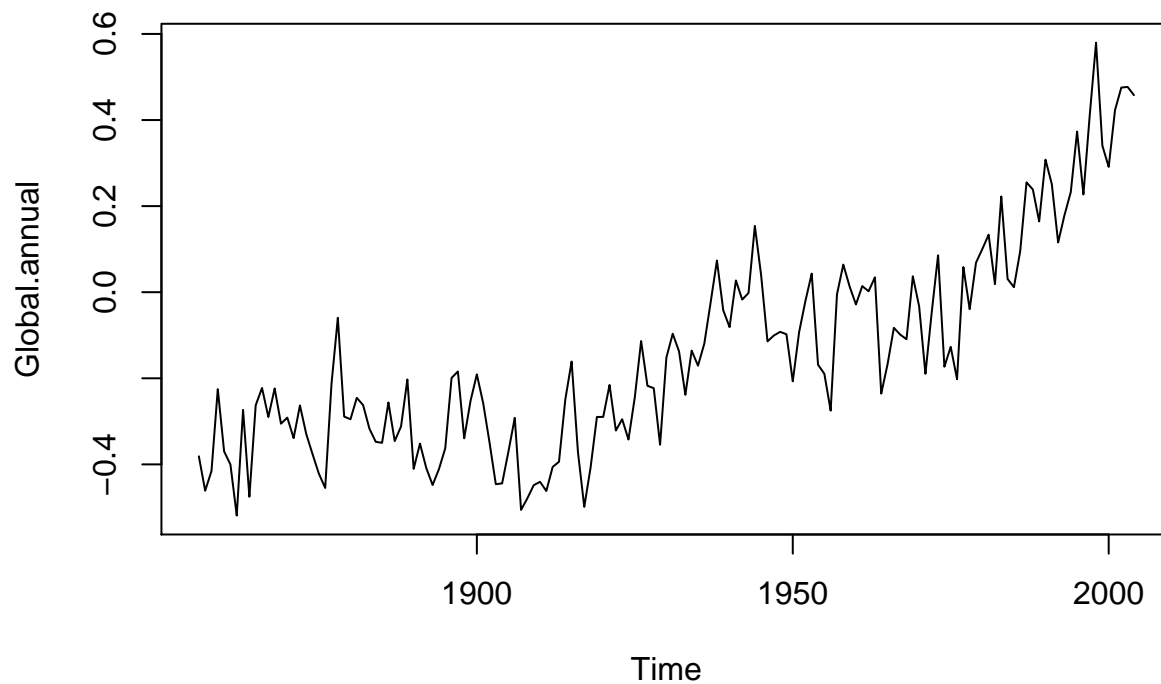


```
## [111] -0.082750000 -0.098416667 -0.109333333 0.037000000 -0.032500000
## [116] -0.189500000 -0.045833333 0.085833333 -0.173083333 -0.127000000
## [121] -0.202250000 0.058416667 -0.039416667 0.068416667 0.100333333
## [126] 0.133583333 0.018666667 0.222416667 0.030666667 0.011666667
## [131] 0.095500000 0.255166667 0.238666667 0.164416667 0.307916667
## [136] 0.251250000 0.115666667 0.178666667 0.232583333 0.373166667
## [141] 0.227000000 0.411000000 0.579666667 0.340500000 0.291083333
## [146] 0.422916667 0.475416667 0.476916667 0.457750000
```

```
plot(GlobalTemperatures.ts)
```



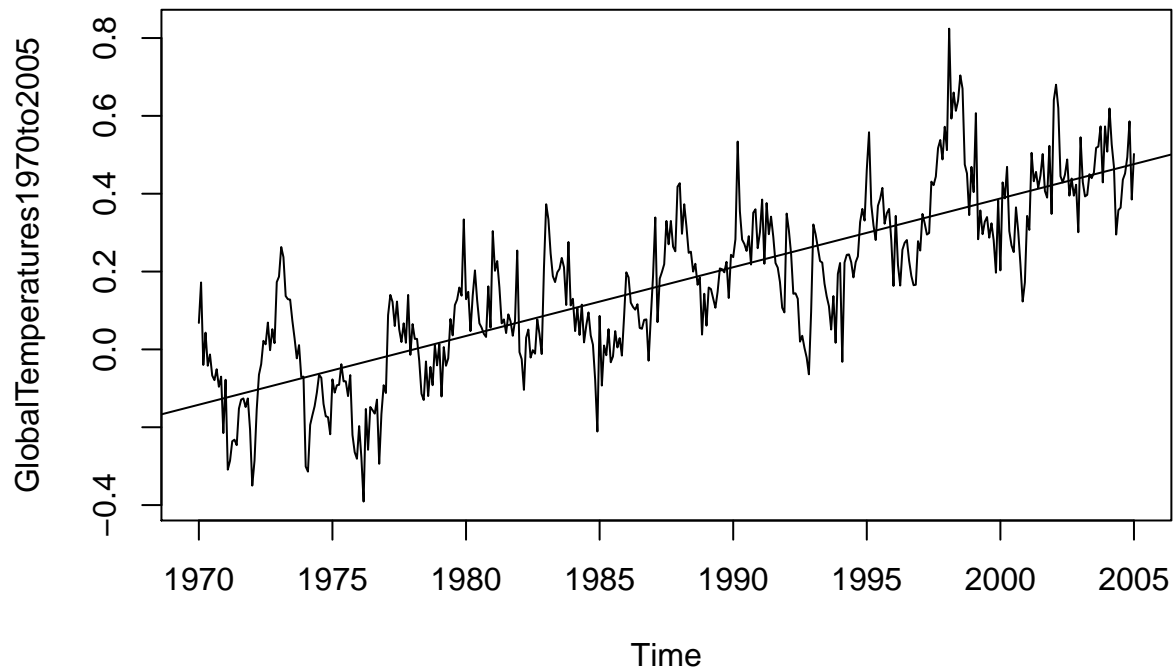
```
plot(Global.annual)
```



```
#
GlobalTemperatures1970to2005 <- window(GlobalTemperatures.ts, start=c(1970, 1), end=c(2005, 12))

## Warning in window.default(x, ...): 'end' value not changed

# subset a time window
?window
GlobalTemperatures1970to2005.time <- time(GlobalTemperatures1970to2005)
# create the vector of times at which a time series was sampled
?time
plot(GlobalTemperatures1970to2005)
abline(reg = lm(GlobalTemperatures1970to2005 ~ GlobalTemperatures1970to2005.time))
```



In statistics, usually, first thing, we compute the ‘mean’ and ‘variance/standard deviation’, which show ‘location’ and ‘dispersion’:

- mean \sim location
- variance \sim dispersion

or

- mean \sim central location
- variance \sim the spread

Transformations

Transform the data (e.g. log or square-root) to stabilize the variance, when:

1. if variance appears to increase with the mean.
 - if standard deviation is directly proportional to the mean (a trend), LOG transform!
 - if variance changes through time, but without a trend, transformation won't help. use a model that accommodates variance change
2. if there's additive seasonal effect, i.e. size of seasonal effect increases with the mean, transform! to make seasonal effect constant.
 - If there's multiplicative seasonal effect, i.e. size of seasonal effect proportional to the mean, LOG transform! to make the seasonal effect additive. (variance gets stable, but error term will still remain unstable)
3. if there are spikes in the data, (skewness), transform to normalize the data distribution.

avoid transformations, except where the transformed has a direct physical interpretation

Sample autocorrelation coefficients

- a series of quantities, that measure the correlation between observations at different distances apart
- show us which probability model to use for that data
- negative correlation? shows high values of x tend to go with low values of y
- zero correlation? shows two variables are independent

serial correlation coefficients (e.g. at lag 1), or autocorrelation coefficients

- measures correlation between successive observations
- serial correlation coefficients at lag k
- if you graph it, it's called 'correlogram'
- ac.f and correlogram is meaningful ONLY IF data is STATIONARY
- ac.f and correlogram is meaningless for non-stationary

how to interpret a Correlogram

- if $r_0 = 1$, r_1 = large value, r_2, r_3 = diminishingly successively smaller values, r_k = almost zero then, TS: one observation 'above' the mean tends to be followed by more observations 'above' the mean, and one observation 'below' the mean tends to be followed by more observations 'below' the mean
- if $r_0 = 1$, r_1 = negative value, r_2 = positive value, r_3 = negative, ... (diminishingly successively smaller values) then, TS: one observation 'above' the mean tends to be followed by more observations 'below' the mean, aka alternating: successive observations on both sides of the overall mean
- if $r_0 \dots r_k$ are all positive and large values, TS is non-stationary, and correlogram is meaningless
- if $r_0 \dots r_k$ oscillate, TS is 'seasonal', e.g. sinusoidal. check at least 3 seasons worth of r_k

Covariance

$$Cov(x, y) = E(xy) - E(x)E(y)$$
$$Var(x + y) = Var(x) + Var(y) + 2Cov(x, y)$$

Basic stochastic models

- residual error series = observed data - fitted values (from the model)
- if our model is good (i.e. captures the deterministic features of the TS)
- then residual TS should be a realization of independent random variables from a probability distribution

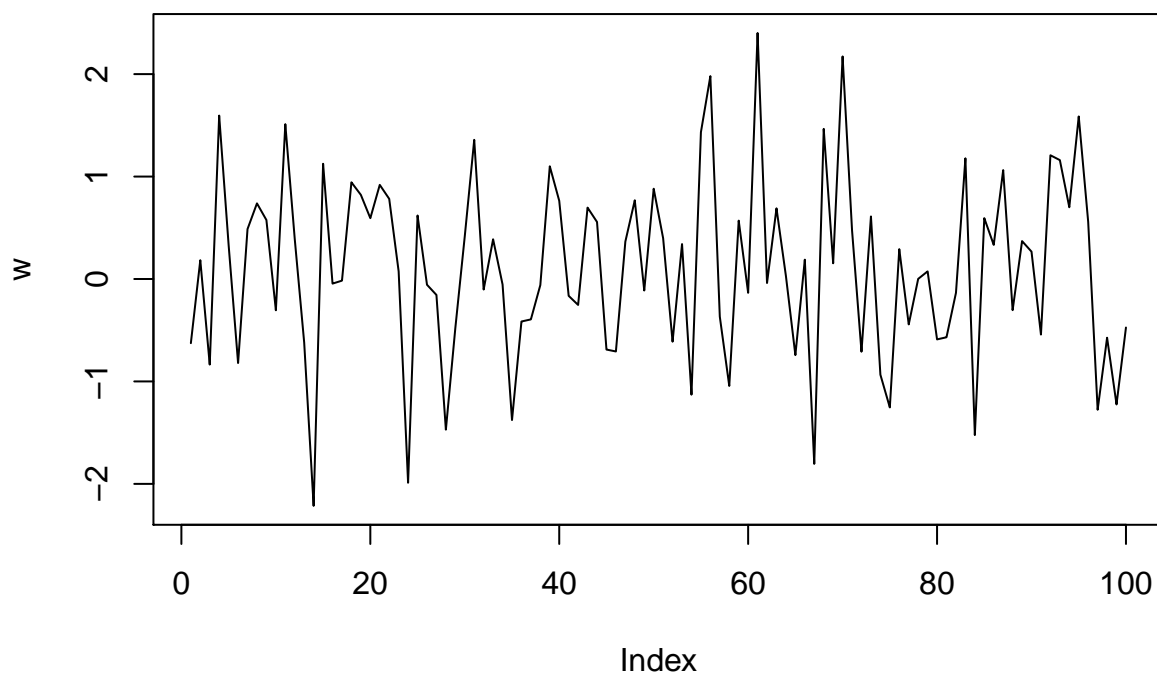
White Noise (WN) [stochastic model for modeling residual TS]

- a TS is discrete WN, if the observations (variables) w_1, \dots, w_t are I.I.D. with mean = zero, and all variables have same variance σ^2 and their pairwise Covariance is zero.
- if, additionally, they are normally distributed, then it's Gaussian WN.
- mean = 0
- $Cov(w_i, w_j) = 0$ for all $i \neq j$

Simulation: to make a synthetic TS (vs. observed TS) Why simulate?

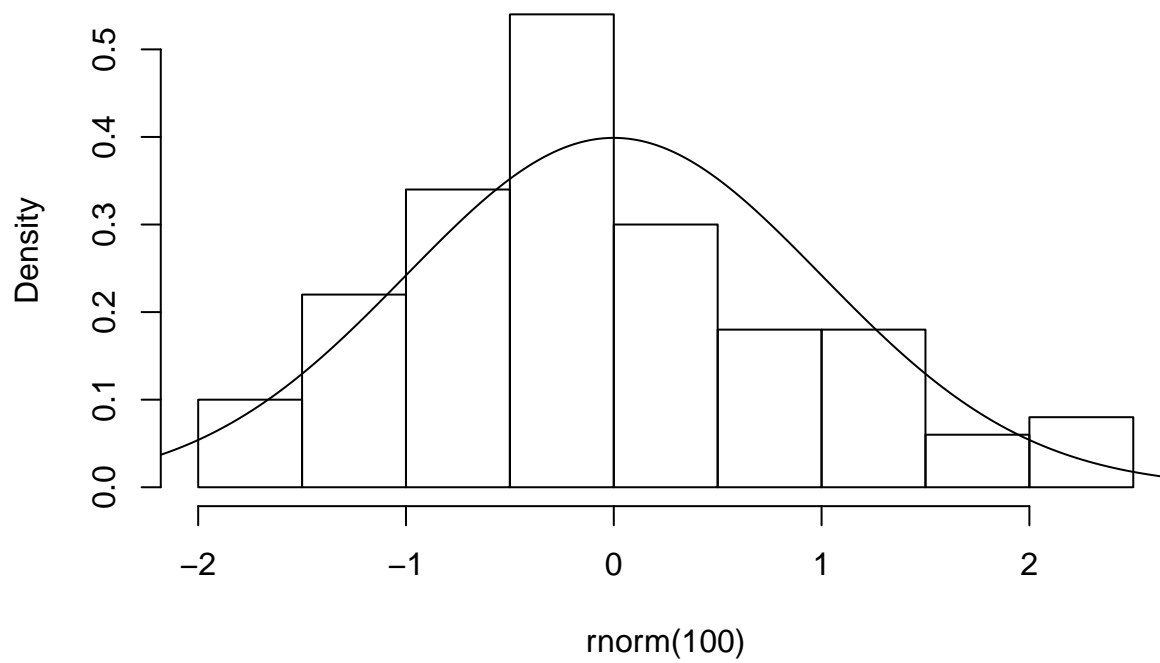
- generate plausible future scenarios
- bootstrapping: constructing confidence intervals for the model parameters

```
# provide a starting point for the random generation function, to make sure the random generation can b  
set.seed(1)  
  
# rnorm simulates 100 random variables that are standard normal and independent  
# which can be used as a Gaussian WN TS of length 100  
w <- rnorm(100)  
  
plot(w, type = "l")
```



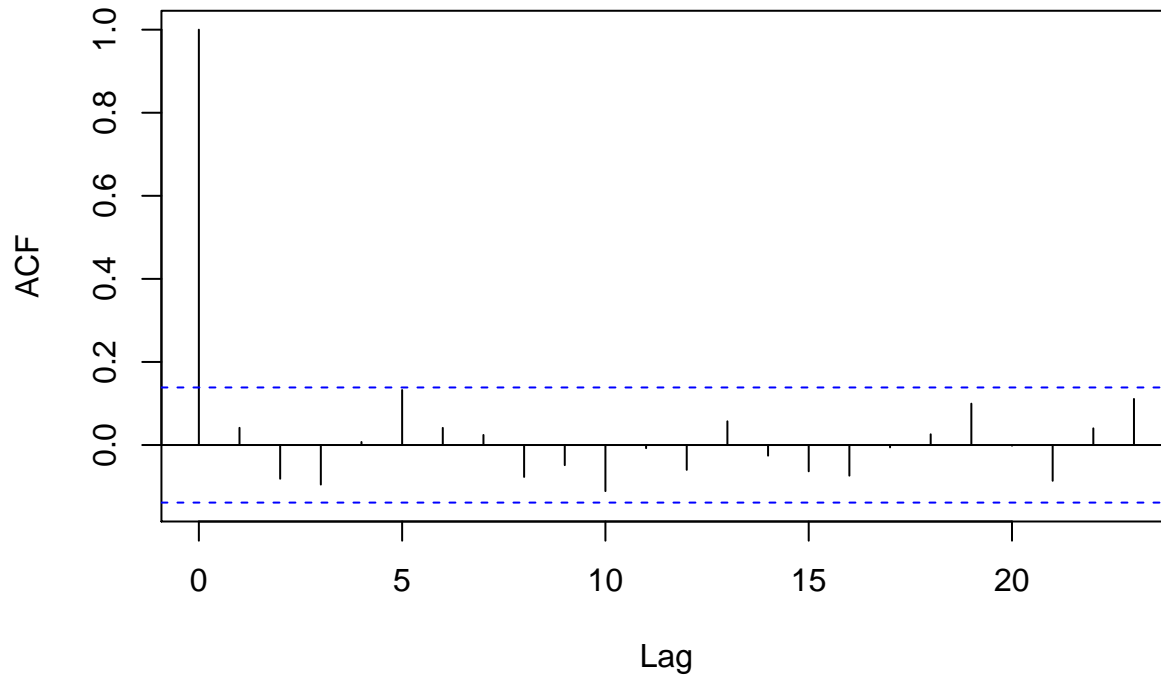
```
# ?dnorm  
# w2 <- dnorm(w, mean = 0, sd = 1, log = FALSE)  
# plot(w2, type = "l")  
  
x <- seq(-3,3, length = 1000)  
hist(rnorm(100), prob = T)  
points(x, dnorm(x), type = "l")
```

Histogram of rnorm(100)



```
# correlogram of WN TS  
set.seed(200)  
acf(rnorm(200))
```

Series rnorm(200)



the correlations (for lag > 1) is (almost because of sampling variation) zero for almost all (95%)

Random Walk (RW) TS [stochastic model for modeling residual TS]

is a good model to fit to data with stochastic trends (not as good as ARIMA though)

TS x_t is a RW, if

$$x_t = x_{t-1} + w_t$$

, where w_t is a WN TS. So, we can back substitute x_{t-1} with $x_{t-2} + w_{t-1}$ and so on, and get:

$$x_t = w_1 + w_2 + w_3 + \dots + w_t$$

(i.e. a finite sum of WN, each with mean = zero and var = σ^2)

- mean = 0
- $Cov(x_t, x_{t+k}) = t\sigma^2 \Rightarrow$ TS is non-stationary because the Covariance depends on time t. variance increases as t increases \Rightarrow RW model is good ONLY for short-term predictions

backshift(lag) operator (B) $Bx_t = x_{t-1}$ $B^n x_t = x_{t-n}$

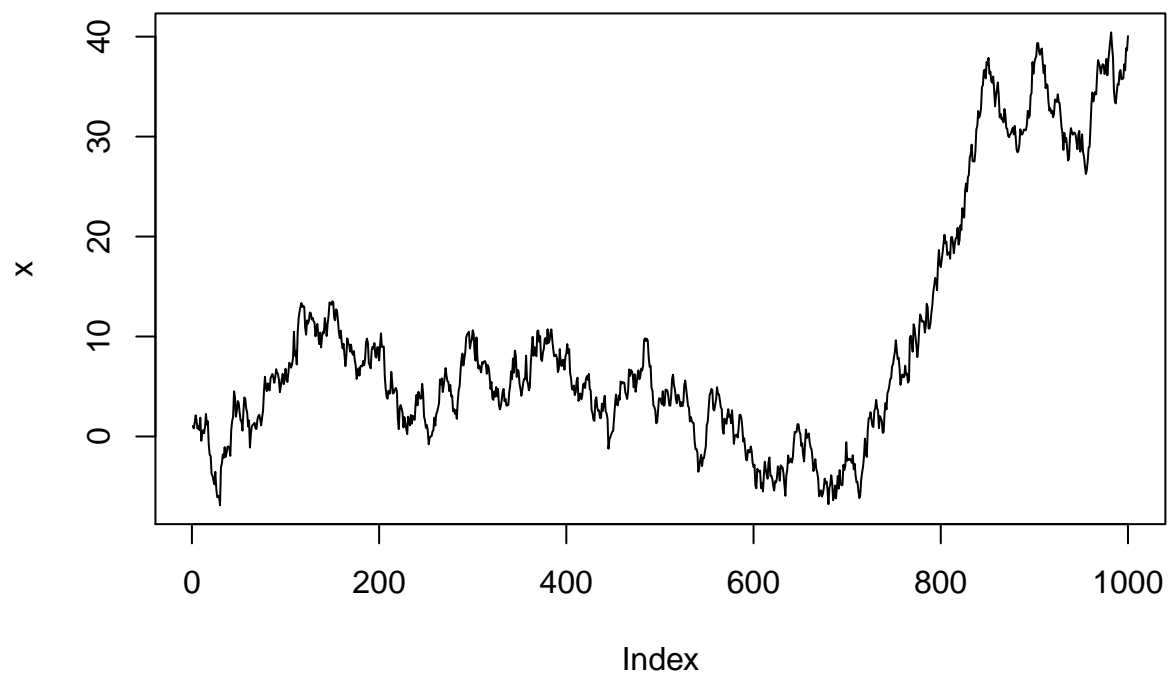
difference operator (to make a stationary TS from a non-stationary TS)

$$\Delta x_t = x_t - x_{t-1}$$


$$\Delta x_t = (1 - B)x_t$$

$$\Delta^n = (1 - B)^n$$

```
# generate WN TS as w
x <- w <- rnorm(1000)
# make a RW TS using backshift and WN TS
for (t in 2:1000) x[t] <- x[t - 1] + w[t]
# plot the RW TS with lines
plot(x, type = "l")
```




```
# draw the RW TS correlogram
acf(x)
```

diagnosis of a RW TS: its correlogram should look like this:

```
?diff
# this correlogram shows a RW TS (as a diagnostic tool)
acf(diff(x))
```

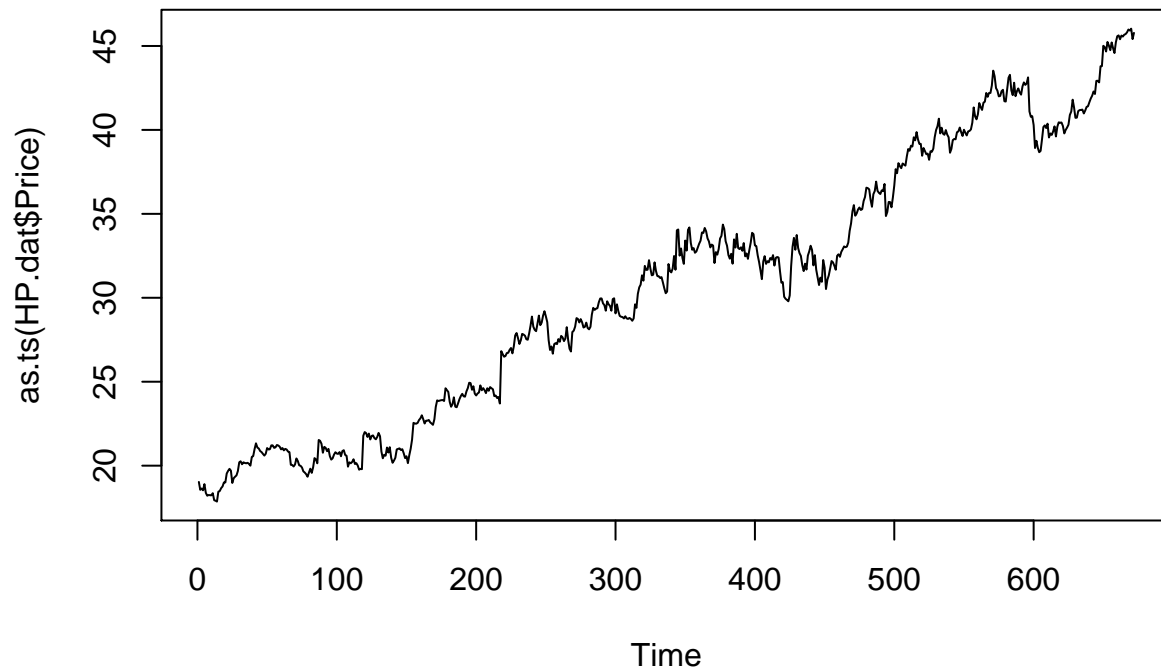


Random Walk with drift TS [stochastic model for modeling residual TS]

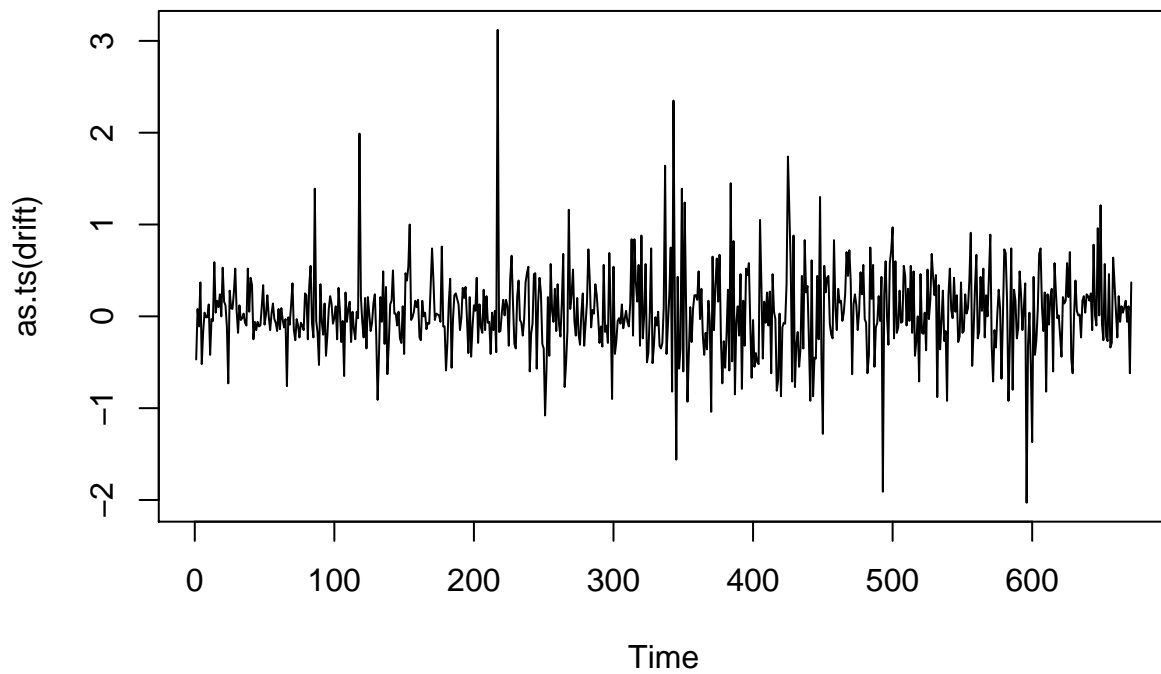
e.g. for stocks, stockholders expect the value of their investment to increase despite the volatility of financial markets. (this increase can be mapped by a drift (upwards) parameter (δ))

$$x_t = x_{t-1} + \delta + w_t$$

```
# HP stock prices
HP.dat <- read.table("http://staff.elena.aut.ac.nz/Paul-Cowpertwait/ts/HP.txt", header = T)
attach(HP.dat)
# plot the stock price TS
plot(as.ts(HP.dat$Price))
```




```
# calculate drift parameter \delta  
drift <- diff(Price)  
plot (as.ts(drift))
```



```
# check correlogram of drift to make sure it's a WN TS  
acf(drift)
```

Series drift



```
# calculate confidence interval
mean(drift) + c(-2, 2) * sd(drift)/sqrt(length(drift))
```

```
## [1] 0.004378275 0.075353468
```

AutoRegressive TS of order p AR(p)[stochastic model for modeling residual TS]

when the model is a regression of x_t on its past terms from the same TS.

$$x_t = \alpha_1 x_{t-1} + \alpha_2 x_{t-2} + \dots + \alpha_p x_{t-p} + w_t$$

where

- w_t is WN TS
- α_i are model parameters
- $\alpha_p \neq 0$ for an AR(p) TS

i.e. value of TS at time t depends not only on t-1 but t-2 and ... t-p (i.e. p steps back)

$$\theta(B)x_t = (1 - \alpha_1 B^1 - \alpha_2 B^2 - \dots - \alpha_p B^p)x_t = w_t$$



how to tell if an AR(p) TS is stationary or non-stationary? $\theta(B) = 1 - \alpha_1 B^1 - \alpha_2 B^2 - \dots - \alpha_p B^p$ is called characteristic equation (CE).

- If all roots of CE for an AR(p) are > 1 , (their absolute value is > 1), AR(p) is stationary.
- If even one of the roots of CE is ≤ 1 in absolute value, then AR(p) is non-stationary.


```
# find roots of a polynomial
polyroot(c(1, 2, 1))
```

```
## [1] -1-0i -1+0i
```


```
rho <- function(k, alpha) alpha^k
layout(1:2)
plot(0:10, rho(0:10, 0.7), type = "b")
plot(0:10, rho(0:10, -0.7), type = "b")
```



```
set.seed(1)
x <- w <- rnorm(100)
for (t in 2:100) x[t] <- 0.7 * x[t - 1] + w[t]
plot(x, type = "l")
acf(x)
```




Series x



```
pacf(x)
```

Series x



```
#### AR(p) TS model fitting in R
```

```
# fit an AutoRegression AR(p) model
```

```
x.ar <- ar(x, method = "mle")
```

```
x.ar$order
```

```
## [1] 1
```

```
# parameter estimate (alpha bar)
```

```
x.ar$ar
```

```
## [1] 0.6009459
```

```
x.ar$asy.var
```

```
##           [,1]
## [1,] 0.006443494
```

```
x.ar$ar + c(-2, 2) * sqrt(x.ar$asy.var)
```

```
## [1] 0.4404031 0.7614886
```

```
# str(x.ar)
```

ARMA(p,q) models: AR(p) + MA(q) combined together: AutoRegressive Moving Average

- AR(p):

$$x_t = \alpha_1 x_{t-1} + \alpha_2 x_{t-2} + \dots + \alpha_p x_{t-p} + w_t$$

- MA(q):

$$x_t = w_t + \beta_1 w_{t-1} + \dots + \beta_q w_{t-q}$$

- ARMA(p,q):

$$x_t = \alpha_1 x_{t-1} + \alpha_2 x_{t-2} + \dots + \alpha_p x_{t-p} + w_t + \beta_1 w_{t-1} + \dots + \beta_q w_{t-q}$$

$$\theta_p(B)x_t = \phi_q(B)w_t$$

- ARMA(p,q) TS is STATIONARY if roots of $\theta > 1$ in absolute value
- ARMA(p,q) TS is INVERTIBLE if roots of $\phi > 1$ in absolute value
- ARMA(p, 0) = AR(p)
- ARMA(0,q) = MA(q)
- ARMA(p,q) is better (parameter efficient) than either a single AR(p) or MA(q)

```
set.seed(1)
# simulate an ARMA data with alpha = -0.6 and beta = 0.5
x <- arima.sim(n = 10000, list(ar = -0.6, ma = 0.5))
# fit an ARMA(1,1) model to the simulated data using c(p,0,q)
coef(arima(x, order = c(1, 0, 1)))
```

```
##           ar1           ma1      intercept
## -0.596966371  0.502703368 -0.006571345
```

```
x.ma <- arima(x.ts, order = c(0, 0, 1))
x.ar <- arima(x.ts, order = c(1, 0, 0))
x.arma <- arima(x.ts, order = c(1, 0, 1))
AIC(x.ma)
```

Now, let's compare AR(p), MA(q), and ARMA(p,q)

```
## [1] -3.526895
```

```
AIC(x.ar)
```

```
## [1] -37.40417
```

```
AIC(x.arma)
```


```
## [1] -42.27357
```

```
# the smaller the AIC or BIC, the better the fit => ARMA(1,1) provides the best fit  
x.arma
```

```
##  
## Call:  
## arima(x = x.ts, order = c(1, 0, 1))  
##  
## Coefficients:  
##          ar1      ma1  intercept  
##         0.8925  0.5319      2.9597  
## s.e.    0.0759  0.2021      0.2435  
##  
## sigma^2 estimated as 0.01505:  log likelihood = 25.14,  aic = -42.27
```

```
# correlogram of residuals of fitted ARMA(1,1) model  
acf(resid(x.arma))
```

Series resid(x.arma)



```
#### example: electricity production series
```




```

CBE <- read.table("http://staff.elena.aut.ac.nz/Paul-Cowpertwait/ts/cbe.dat", header = T)
Elec.ts <- ts(CBE[, 3], start = 1958, freq = 12)
Time <- 1:length(Elec.ts)
Imth <- cycle(Elec.ts)
Elec.lm <- lm(log(Elec.ts) ~ Time + I(Time^2) + factor(Imth))
acf(resid(Elec.lm))

```

Series resid(Elec.lm)



```

# find the best p and q for ARMA(p,q)
best.order <- c(0, 0, 0)
best.aic <- Inf
for (i in 0:2) for (j in 0:2) {
  fit.aic <- AIC(arima(resid(Elec.lm), order = c(i, 0,j)))
  if (fit.aic < best.aic) {
    best.order <- c(i, 0, j)
    best.arma <- arima(resid(Elec.lm), order = best.order)
    best.aic <- fit.aic
  }
}
# what's the best order found:
best.order

```


```
## [1] 2 0 0
```

```

# correlogram
acf(resid(best.arma))

```


Series resid(best.arma)



prediction

something

```
new.time <- seq(length(Elec.ts), length = 36)
new.data <- data.frame(Time = new.time, lmth = rep(1:12,3))
# predict using the fitted regression model of class lm
predict.lm <- predict(Elec.lm, new.data)
# predict using the fitted ARMA model, number of timesteps ahead to predict; here = 36 ~ predict for th
predict.arma <- predict(best.arma, n.ahead = 36)
# convert the predicted values to TS object
elec.pred <- ts(exp(predict.lm + predict.arma$pred), start = 1991, freq = 12)
ts.plot(cbind(Elec.ts, elec.pred), lty = 1:2)
```



— ### how to tell if a correlogram is of AR(1), MA(1) or WN

AR(1): smooth decay AR(1) with negative correlation: smooth decay, alternating positive and negative

MA(1): lag 0 = 1, lag 1 = beta_1 WN: lag 0 = 1, the rest are almost zero.

Regression

Time series regression vs. regular regression? TS regression is different because its residual forms a TS, and is serially correlated. (if this correlation is positive => estimated s.e. of parameter estimates are less than their true values => erroneously low p-values) how to fix this? and get a proper estimated s.e.? use GLS (generalized least squares) to calculate estimated s.e.


To do regression, we need to make the TS stationary, if it's non-stationary. e.g. if TS is $x_t = \alpha_0 + \alpha_1 t + z_t$, which is non-stationary, because it depends on t , then,

$$\Delta x_t = x_t - x_{t-1} = \alpha_1 + z_t - z_{t-1}$$

i.e. first order differencing, removes the trend, and makes it stationary. (assuming that z_t is stationary error series)

If the underlying trend is a polynomial of order m , then we need to do m^{th} -order differencing to remove its trend.

```
set.seed(1)
# Gaussian WN
z <- w <- rnorm(100, sd = 20)
# noise = AR(1) with alpha 0.8 and WN
for (t in 2:100) z[t] <- 0.8 * z[t - 1] + w[t]
Time <- 1:100
# TS with a straight line trend and AR(1) residual error noise
x <- 50 + 3 * Time + z
plot(x, xlab = "time", type = "l")
```



```
# fit a linear regression model
x.lm <- lm(x ~ Time)
# look at the coefficients of the estimated linear regression model
coef(x.lm)
```

```
## (Intercept)      Time
##  58.551218    3.063275
```

```
# calculate s.e. (standard errors): will be under-estimated because of serial correlation
sqrt(diag(vcov(x.lm)))
```

```
## (Intercept)      Time
##  4.88006278  0.08389621
```


```
# summary() will give erroneous info, because of the under-estimated s.e. => incorrect p-values
summary(x.lm)
```

```
##
## Call:
## lm(formula = x ~ Time)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -42.30 -16.44  -0.54   13.66   57.00
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  58.5512    4.8801   12.00  <2e-16 ***
## Time         3.0633     0.0839   36.51  <2e-16 ***
```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 24.22 on 98 degrees of freedom
## Multiple R-squared:  0.9315, Adjusted R-squared:  0.9308
## F-statistic: 1333 on 1 and 98 DF,  p-value: < 2.2e-16


# check diagnostic plots
acf(residuals(x.lm))
```

Series residuals(x.lm)



```
pacf(residuals(x.lm))
```

Series residuals(x.lm)



another example:

```
Global <- scan("http://staff.elena.aut.ac.nz/Paul-Cowpertwait/ts/global.dat")
Global.ts <- ts(Global, st = c(1856, 1), end = c(2005, 12), fr = 12)
temp <- window(Global.ts, start = 1970)
temp.lm <- lm(temp ~ time(temp))
coef(temp.lm)
```


```
## (Intercept)  time(temp)
## -34.920409    0.017654
```

```
confint(temp.lm)
```

```
##                2.5 %        97.5 %
## (Intercept) -37.21001248 -32.63080554
## time(temp)   0.01650228  0.01880572
```

```
acf(resid(lm(temp ~ time(temp))))
```

Series resid(lm(temp ~ time(temp)))



GLS (generalized least squares)

- GLS is a fitting procedure
- to estimate the s.e. of the regression model parameters
- better than the OLS we used before

```
library(nlme)
x.gls <- gls(x ~ Time, cor = corAR1(0.8))
summary(x.gls)

## Generalized least squares fit by REML
## Model: x ~ Time
## Data: NULL
##      AIC      BIC    logLik
## 862.8866 873.2265 -427.4433
##
## Correlation Structure: AR(1)
## Formula: ~1
## Parameter estimate(s):
##      Phi
## 0.7161368
##
## Coefficients:
##              Value Std.Error   t-value p-value
## (Intercept) 58.23302 11.924568   4.883449      0
## Time         3.04225  0.202445  15.027538      0
```

```
##
## Correlation:
## (Intr)
## Time -0.857
##
## Standardized residuals:
##      Min      Q1      Med      Q3      Max
## -1.6171329 -0.6195428  0.0353972  0.5836326  2.3184155
##
## Residual standard error: 25.58595
## Degrees of freedom: 100 total; 98 residual
```

```
temp.gls <- gls(temp ~ time(temp), cor = corAR1(0.7))
summary(temp.gls)
```

```
## Generalized least squares fit by REML
## Model: temp ~ time(temp)
## Data: NULL
##      AIC      BIC    logLik
## -838.4383 -822.1832 423.2192
##
## Correlation Structure: AR(1)
## Formula: ~1
## Parameter estimate(s):
##      Phi
## 0.7206941
##
## Coefficients:
##      Value Std.Error   t-value p-value
## (Intercept) -34.15115  2.8850347 -11.83734      0
## time(temp)   0.01727  0.0014512  11.89822      0
##
## Correlation:
##      (Intr)
## time(temp) -1
##
## Standardized residuals:
##      Min      Q1      Med      Q3      Max
## -2.81994762 -0.65516626 -0.02914169  0.60700264  3.68710752
##
## Residual standard error: 0.1285624
## Degrees of freedom: 432 total; 430 residual
```

```
# confidence interval
confint(temp.gls)
```

```
##      2.5 %      97.5 %
## (Intercept) -39.80571681 -28.49658850
## time(temp)   0.01442274  0.02011148
```

```
# note that the confidence interval does not include zero => estimates are statistically significant =>
```


Linear models with seasonality

treat the seasonal term as a factor! (e.g. 12 factors for monthly data TS) and do GLS or OLS like before.

```
Seas <- cycle(temp)
Time <- time(temp)
# factors for seasonality. Use 0 in model to make sure the model won't have an intercept, otherwise one
temp.lm <- lm(temp ~ 0 + Time + factor(Seas))
summary(temp.lm)
```

```
##
## Call:
## lm(formula = temp ~ 0 + Time + factor(Seas))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.37411 -0.08045 -0.00143  0.07810  0.43079
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## Time              1.771e-02  5.849e-04   30.27  <2e-16 ***
## factor(Seas)1    -3.500e+01  1.163e+00  -30.10  <2e-16 ***
## factor(Seas)2    -3.499e+01  1.163e+00  -30.09  <2e-16 ***
## factor(Seas)3    -3.501e+01  1.163e+00  -30.11  <2e-16 ***
## factor(Seas)4    -3.501e+01  1.163e+00  -30.11  <2e-16 ***
## factor(Seas)5    -3.503e+01  1.163e+00  -30.12  <2e-16 ***
## factor(Seas)6    -3.503e+01  1.163e+00  -30.12  <2e-16 ***
## factor(Seas)7    -3.503e+01  1.163e+00  -30.12  <2e-16 ***
## factor(Seas)8    -3.502e+01  1.163e+00  -30.11  <2e-16 ***
## factor(Seas)9    -3.504e+01  1.163e+00  -30.12  <2e-16 ***
## factor(Seas)10   -3.505e+01  1.163e+00  -30.14  <2e-16 ***
## factor(Seas)11   -3.507e+01  1.163e+00  -30.14  <2e-16 ***
## factor(Seas)12   -3.505e+01  1.163e+00  -30.13  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1263 on 419 degrees of freedom
## Multiple R-squared:  0.8072, Adjusted R-squared:  0.8012
## F-statistic: 134.9 on 13 and 419 DF, p-value: < 2.2e-16
```

```
coef(temp.lm)
```

```
##           Time factor(Seas)1 factor(Seas)2 factor(Seas)3 factor(Seas)4
##    0.01770758  -34.99726483  -34.98801824  -35.01002165  -35.01227506
## factor(Seas)5 factor(Seas)6 factor(Seas)7 factor(Seas)8 factor(Seas)9
##   -35.03369513  -35.02505965  -35.02689640  -35.02476092  -35.03831988
## factor(Seas)10 factor(Seas)11 factor(Seas)12
##   -35.05248996  -35.06557670  -35.04871900
```

```
# predict two years ahead
new.t <- seq(2006, len = 2 * 12, by = 1/12)
new.dat <- data.frame(Time = new.t, Seas = rep(1:12, 2))
predict(temp.lm, new.dat)[1:24]
```

```
##          1          2          3          4          5          6          7
## 0.5241458 0.5348681 0.5143403 0.5135625 0.4936181 0.5037292 0.5033681
##          8          9         10         11         12         13         14
## 0.5069792 0.4948958 0.4822014 0.4705903 0.4889236 0.5418534 0.5525756
##         15         16         17         18         19         20         21
## 0.5320479 0.5312701 0.5113256 0.5214367 0.5210756 0.5246867 0.5126034
##         22         23         24
## 0.4999090 0.4882979 0.5066312
```

Harmonic seasonal models

instead of using one parameter per season, smooth over the seasons => more parameter-efficient

A sine wave with frequency f (cycles per sampling interval), amplitude A , and phase shift ϕ can be expressed as:

$$A \sin(2\pi ft + \phi) = \alpha_s \sin(2\pi ft) + \alpha_c \cos(2\pi ft)$$

For a time series $\{x_t\}$ with s seasons there are $\lfloor s/2 \rfloor$ possible cycles. e.g. for monthly data, $s = 12$, there are 6 possible cycles.


The model will be:

$$x_t = m_t + \sum_{i=1}^{\lfloor s/2 \rfloor} (s_i \sin(2\pi it/s) + c_i \cos(2\pi it/s)) + z_t$$


where:

- m_t is trend
- s_i and c_i unknown parameters


```
TIME <- seq(1, 12, len = 1000)
plot(TIME, sin(2 * pi * TIME/12), type = "l")
```



```
plot(TIME, sin(2 * pi * TIME/12) +  
      0.2 * sin(2 * pi * 2 * TIME/12) +  
      0.1 * sin(2 * pi * 4 * TIME/12) +  
      0.1 * cos(2 * pi * 4 * TIME/12),  
      type = "l")
```



```
set.seed(1)
TIME <- 1:(10 * 12)
w <- rnorm(10 * 12, sd = 0.5)
Trend <- 0.1 + 0.005 * TIME + 0.001 * TIME^2
Seasonal <- sin(2*pi*TIME/12) +
  0.2*sin(2*pi*2*TIME/12) +
  0.1*sin(2*pi*4*TIME/12) +
  0.1*cos(2*pi*4*TIME/12)
x <- Trend + Seasonal + w
plot(x, type = "l")
```



```
SIN <- COS <- matrix(nr = length(TIME), nc = 6)
for (i in 1:6) {
  COS[, i] <- cos(2 * pi * i * TIME/12)
  SIN[, i] <- sin(2 * pi * i * TIME/12)
}
```

```
x.lm1 <- lm(x ~ TIME + I(TIME^2) +
  COS[, 1] + SIN[, 1] +
  COS[, 2] + SIN[, 2] +
  COS[, 3] + SIN[, 3] +
  COS[, 4] + SIN[, 4] +
  COS[, 5] + SIN[, 5] +
  COS[, 6] + SIN[, 6])
coef(x.lm1)/sqrt(diag(vcov(x.lm1)))
```

```
## (Intercept)      TIME    I(TIME^2)    COS[, 1]    SIN[, 1]    COS[, 2]
##  1.2390758    1.1251009    25.9327225    0.3277392    15.4421534   -0.5146340
##      SIN[, 2]    COS[, 3]    SIN[, 3]    COS[, 4]    SIN[, 4]    COS[, 5]
##  3.4468638    0.2317971   -0.7027374    0.2276281    1.0525106   -1.1501059
##      SIN[, 5]    COS[, 6]    SIN[, 6]
##  0.8573584   -0.3100116    0.3823830
```

```
x.lm2 <- lm(x ~ I(TIME^2) + SIN[, 1] + SIN[, 2])
coef(x.lm2)/sqrt(diag(vcov(x.lm2)))
```

```
## (Intercept)    I(TIME^2)    SIN[, 1]    SIN[, 2]
##    4.627904   111.143693    15.786425    3.494708
```

```
coef(x.lm2)
```

```
## (Intercept)    I(TIME^2)    SIN[, 1]    SIN[, 2]  
## 0.280403915 0.001036409 0.900206804 0.198862866
```

```
AIC(x.lm1)
```

```
## [1] 165.3527
```

```
AIC(x.lm2)
```

```
## [1] 149.7256
```

```
AIC(lm(x ~ TIME + I(TIME^2) +  
      SIN[,1] +  
      SIN[,2] +  
      SIN[,4] +  
      COS[,4]))
```

```
## [1] 153.0309
```

```
step(x.lm1)
```

```
## Start: AIC=-177.19  
## x ~ TIME + I(TIME^2) + COS[, 1] + SIN[, 1] + COS[, 2] + SIN[,  
##      2] + COS[, 3] + SIN[, 3] + COS[, 4] + SIN[, 4] + COS[, 5] +  
##      SIN[, 5] + COS[, 6] + SIN[, 6]  
##  
##           Df Sum of Sq    RSS    AIC  
## - COS[, 4]  1      0.011 21.357 -179.133  
## - COS[, 3]  1      0.011 21.357 -179.131  
## - COS[, 6]  1      0.020 21.366 -179.083  
## - COS[, 1]  1      0.022 21.368 -179.070  
## - SIN[, 6]  1      0.030 21.376 -179.026  
## - COS[, 2]  1      0.054 21.400 -178.890  
## - SIN[, 3]  1      0.100 21.447 -178.630  
## - SIN[, 5]  1      0.149 21.496 -178.355  
## - SIN[, 4]  1      0.225 21.572 -177.933  
## - TIME      1      0.257 21.604 -177.755  
## - COS[, 5]  1      0.269 21.615 -177.690  
## <none>                21.346 -177.193  
## - SIN[, 2]  1      2.415 23.762 -166.329  
## - SIN[, 1]  1     48.479 69.825  -36.979  
## - I(TIME^2)  1    136.720 158.067   61.063  
##  
## Step: AIC=-179.13  
## x ~ TIME + I(TIME^2) + COS[, 1] + SIN[, 1] + COS[, 2] + SIN[,  
##      2] + COS[, 3] + SIN[, 3] + SIN[, 4] + COS[, 5] + SIN[, 5] +  
##      COS[, 6] + SIN[, 6]  
##
```

```

##           Df Sum of Sq      RSS       AIC
## - COS[, 3]    1      0.011   21.368 -181.072
## - COS[, 6]    1      0.019   21.376 -181.029
## - COS[, 1]    1      0.022   21.379 -181.011
## - SIN[, 6]    1      0.032   21.389 -180.952
## - COS[, 2]    1      0.054   21.411 -180.830
## - SIN[, 3]    1      0.100   21.457 -180.572
## - SIN[, 5]    1      0.150   21.507 -180.292
## - SIN[, 4]    1      0.225   21.582 -179.876
## - TIME        1      0.258   21.615 -179.694
## - COS[, 5]    1      0.270   21.627 -179.625
## <none>                21.357 -179.133
## - SIN[, 2]    1      2.419   23.776 -168.260
## - SIN[, 1]    1     48.480   69.837  -38.959
## - I(TIME^2)    1    136.726 158.083   59.076
##
## Step:  AIC=-181.07
## x ~ TIME + I(TIME^2) + COS[, 1] + SIN[, 1] + COS[, 2] + SIN[,
##      2] + SIN[, 3] + SIN[, 4] + COS[, 5] + SIN[, 5] + COS[, 6] +
##      SIN[, 6]
##
##           Df Sum of Sq      RSS       AIC
## - COS[, 6]    1      0.019   21.387 -182.964
## - COS[, 1]    1      0.022   21.390 -182.949
## - SIN[, 6]    1      0.031   21.399 -182.898
## - COS[, 2]    1      0.054   21.422 -182.769
## - SIN[, 3]    1      0.100   21.468 -182.510
## - SIN[, 5]    1      0.150   21.518 -182.233
## - SIN[, 4]    1      0.225   21.593 -181.814
## - TIME        1      0.258   21.626 -181.630
## - COS[, 5]    1      0.270   21.638 -181.567
## <none>                21.368 -181.072
## - SIN[, 2]    1      2.417   23.785 -170.211
## - SIN[, 1]    1     48.483   69.851  -40.935
## - I(TIME^2)    1    136.725 158.093   57.083
##
## Step:  AIC=-182.96
## x ~ TIME + I(TIME^2) + COS[, 1] + SIN[, 1] + COS[, 2] + SIN[,
##      2] + SIN[, 3] + SIN[, 4] + COS[, 5] + SIN[, 5] + SIN[, 6]
##
##           Df Sum of Sq      RSS       AIC
## - COS[, 1]    1      0.022   21.409 -184.842
## - COS[, 2]    1      0.057   21.444 -184.645
## - SIN[, 6]    1      0.080   21.467 -184.517
## - SIN[, 3]    1      0.097   21.484 -184.420
## - SIN[, 5]    1      0.160   21.547 -184.069
## - SIN[, 4]    1      0.223   21.610 -183.721
## - TIME        1      0.257   21.644 -183.530
## - COS[, 5]    1      0.283   21.670 -183.387
## <none>                21.387 -182.964
## - SIN[, 2]    1      2.455   23.842 -171.923
## - SIN[, 1]    1     48.475   69.862  -42.916
## - I(TIME^2)    1    136.801 158.188   55.155
##

```

```

## Step: AIC=-184.84
## x ~ TIME + I(TIME^2) + SIN[, 1] + COS[, 2] + SIN[, 2] + SIN[,
##      3] + SIN[, 4] + COS[, 5] + SIN[, 5] + SIN[, 6]
##
##           Df Sum of Sq      RSS      AIC
## - COS[, 2]  1      0.057  21.466 -186.523
## - SIN[, 6]  1      0.080  21.489 -186.396
## - SIN[, 3]  1      0.097  21.506 -186.298
## - SIN[, 5]  1      0.160  21.569 -185.947
## - SIN[, 4]  1      0.223  21.632 -185.599
## - TIME      1      0.257  21.666 -185.412
## - COS[, 5]  1      0.283  21.692 -185.266
## <none>                      21.409 -184.842
## - SIN[, 2]  1      2.456  23.865 -173.812
## - SIN[, 1]  1     48.478  69.887  -44.873
## - I(TIME^2) 1    136.844 158.253   53.205
##
## Step: AIC=-186.52
## x ~ TIME + I(TIME^2) + SIN[, 1] + SIN[, 2] + SIN[, 3] + SIN[,
##      4] + COS[, 5] + SIN[, 5] + SIN[, 6]
##
##           Df Sum of Sq      RSS      AIC
## - SIN[, 6]  1      0.070  21.536 -188.131
## - SIN[, 3]  1      0.098  21.564 -187.976
## - SIN[, 5]  1      0.158  21.624 -187.643
## - SIN[, 4]  1      0.223  21.689 -187.281
## - TIME      1      0.256  21.722 -187.100
## - COS[, 5]  1      0.280  21.746 -186.966
## <none>                      21.466 -186.523
## - SIN[, 2]  1      2.449  23.915 -175.560
## - SIN[, 1]  1     48.475  69.941  -46.781
## - I(TIME^2) 1    136.825 158.290   51.233
##
## Step: AIC=-188.13
## x ~ TIME + I(TIME^2) + SIN[, 1] + SIN[, 2] + SIN[, 3] + SIN[,
##      4] + COS[, 5] + SIN[, 5]
##
##           Df Sum of Sq      RSS      AIC
## - SIN[, 3]  1      0.109  21.645 -189.525
## - SIN[, 5]  1      0.132  21.668 -189.400
## - SIN[, 4]  1      0.232  21.768 -188.847
## - COS[, 5]  1      0.247  21.783 -188.761
## - TIME      1      0.258  21.794 -188.702
## <none>                      21.536 -188.131
## - SIN[, 2]  1      2.386  23.923 -177.520
## - SIN[, 1]  1     48.501  70.037  -48.616
## - I(TIME^2) 1    136.755 158.291   49.233
##
## Step: AIC=-189.53
## x ~ TIME + I(TIME^2) + SIN[, 1] + SIN[, 2] + SIN[, 4] + COS[,
##      5] + SIN[, 5]
##
##           Df Sum of Sq      RSS      AIC
## - SIN[, 5]  1      0.132  21.777 -190.798

```



```

## - SIN[, 4] 1 0.232 21.877 -190.247
## - COS[, 5] 1 0.247 21.893 -190.161
## - TIME 1 0.260 21.905 -190.093
## <none> 21.645 -189.525
## - SIN[, 2] 1 2.387 24.032 -178.971
## - SIN[, 1] 1 48.508 70.154 -50.417
## - I(TIME^2) 1 136.750 158.395 47.312
##
## Step: AIC=-190.8
## x ~ TIME + I(TIME^2) + SIN[, 1] + SIN[, 2] + SIN[, 4] + COS[,
## 5]
##
## Df Sum of Sq RSS AIC
## - SIN[, 4] 1 0.232 22.008 -191.528
## - COS[, 5] 1 0.247 22.024 -191.443
## - TIME 1 0.259 22.036 -191.378
## <none> 21.777 -190.798
## - SIN[, 2] 1 2.387 24.164 -180.317
## - SIN[, 1] 1 48.506 70.283 -52.195
## - I(TIME^2) 1 136.752 158.528 45.413
##
## Step: AIC=-191.53
## x ~ TIME + I(TIME^2) + SIN[, 1] + SIN[, 2] + COS[, 5]
##
## Df Sum of Sq RSS AIC
## - COS[, 5] 1 0.247 22.256 -192.187
## - TIME 1 0.258 22.266 -192.131
## <none> 22.008 -191.528
## - SIN[, 2] 1 2.386 24.395 -181.175
## - SIN[, 1] 1 48.500 70.509 -53.811
## - I(TIME^2) 1 136.756 158.764 43.591
##
## Step: AIC=-192.19
## x ~ TIME + I(TIME^2) + SIN[, 1] + SIN[, 2]
##
## Df Sum of Sq RSS AIC
## - TIME 1 0.255 22.511 -192.820
## <none> 22.256 -192.187
## - SIN[, 2] 1 2.385 24.641 -181.970
## - SIN[, 1] 1 48.490 70.745 -55.409
## - I(TIME^2) 1 136.755 159.011 41.778
##
## Step: AIC=-192.82
## x ~ I(TIME^2) + SIN[, 1] + SIN[, 2]
##
## Df Sum of Sq RSS AIC
## <none> 22.51 -192.82
## - SIN[, 2] 1 2.37 24.88 -182.81
## - SIN[, 1] 1 48.36 70.87 -57.19
## - I(TIME^2) 1 2397.19 2419.70 366.47
##
## Call:
## lm(formula = x ~ I(TIME^2) + SIN[, 1] + SIN[, 2])

```

```
##
## Coefficients:
## (Intercept)      I(TIME^2)      SIN[, 1]      SIN[, 2]
##      0.280404      0.001036      0.900207      0.198863
```

```
SIN <- COS <- matrix(nr = length(temp), nc = 6)
for (i in 1:6) {
  COS[, i] <- cos(2 * pi * i * time(temp))
  SIN[, i] <- sin(2 * pi * i * time(temp))
}
TIME <- (time(temp) - mean(time(temp)))/sd(time(temp))
mean(time(temp))
```

Fit a harmonic model with a quadratic trend is fitted to the temperature series (1970–2005)

```
## [1] 1987.958
```

```
sd(time(temp))
```

```
## [1] 10.40433
```

```
temp.lm1 <- lm(temp ~ TIME + I(TIME^2) +
               COS[,1] + SIN[,1] + COS[,2] + SIN[,2] +
               COS[,3] + SIN[,3] + COS[,4] + SIN[,4] +
               COS[,5] + SIN[,5] + COS[,6] + SIN[,6])
coef(temp.lm1)/sqrt(diag(vcov(temp.lm1)))
```

```
## (Intercept)      TIME      I(TIME^2)      COS[, 1]      SIN[, 1]      COS[, 2]
## 18.2569117 30.2800742  1.2736899  0.7447906  2.3845512  1.2086419
##      SIN[, 2]      COS[, 3]      SIN[, 3]      COS[, 4]      SIN[, 4]      COS[, 5]
##  1.9310853  0.6448116  0.3971873  0.5468025  0.1681753  0.3169782
##      SIN[, 5]      COS[, 6]      SIN[, 6]
##  0.3504607 -0.4498835 -0.6216650
```

```
temp.lm2 <- lm(temp ~ TIME + SIN[, 1] + SIN[, 2])
coef(temp.lm2)
```

```
## (Intercept)      TIME      SIN[, 1]      SIN[, 2]
## 0.17501157 0.18409618 0.02041803 0.01615374
```

```
AIC(temp.lm)
```

```
## [1] -546.9577
```


```
AIC(temp.lm1)
```

```
## [1] -545.0491
```

```
AIC(temp.lm2)
```


```
## [1] -561.1779
```

```
plot(time(temp), resid(temp.lm2), type = "l")  
abline(0, 0, col = "red")
```




```
acf(resid(temp.lm2))
```

Series resid(temp.lm2)



```
pacf(resid(temp.lm2))
```

Series resid(temp.lm2)



```
res.ar <- ar(resid(temp.lm2), method = "mle")
res.ar$ar
```

```
## [1] 0.4938189 0.3071598
```

```
sd(res.ar$res[-(1:2)])
```

```
## [1] 0.08373324
```

```
acf(res.ar$res[-(1:2)])
```

